

Exploratory Data Analysis Notes

Coursera Course by John Hopkins University

INSTRUCTORS: Dr. Jeff Leek, Dr. Roger D. Peng, Dr. Brian Caffo

Contents

Intro	3
Instructor's Note	3
Introduction	3
Exploratory Data Analysis with R Book	4
The Art of Data Science	4
Installing R on...	4
Windows	4
Mac	4
Installing R Studio (Mac)	4
Setting Your Working Directory on...	4
Windows, Mac & Linux	4
Lesson 1: Graphs	5
Principles of Analytic Graphics	5
Lesson with <code>swirl()</code> : Principles of Analytic Graphics	9
Exploratory Graphs	9
Five-number Summary	10
Box plots	10
Histograms	12
Density plot	16
Bar plot	18
Multiple Box plots	19
Multiple Histograms	20
Scatter plot	21
Lesson with <code>swirl()</code> : Exploratory Graphs	24
Lesson 2: Plotting	25
Plotting Systems in R	25
Lesson with <code>swirl()</code> : Plotting Systems	28
Base Plotting System	28
Base Histogram	29
Base Scatter plot	30
Base Box plot	31
Some Important Base Graphics Parameters	32
Base Plotting Functions	33

Multiple Base Plots	37
Summary	39
Base Plotting Demonstration	39
Lesson with <code>swirl()</code> : Base Plotting System	43
Lesson 3: Graphics Devices	45
Graphics Devices in R	45
Lesson with <code>swirl()</code> : Graphics Devices in R	45
Quiz 1 Scribbles	45
Course Project 1	45
Lesson 4: Lattice Plotting	45
Lattice Plotting System (part 1)	45
Lattice Plotting System (part 2)	45
Lesson with <code>swirl()</code> : Lattice Plotting System	45
Lesson with <code>swirl()</code> : Working with Colors	45
Lesson 5: ggplot2 <3	46
Part 1	46
Part 2	46
Lesson with <code>swirl()</code> : GGPlot2 Part 1	46
Part 3	46
Part 4	46
Lesson with <code>swirl()</code> : GGPlot2 Part 2	46
Part 5	46
Lesson with <code>swirl()</code> : GGPlot2 Extras	46
Quiz 2 Scribbles	46
Lesson 6: Hierarchical Clustering	46
Part 1	46
Part 2	46
Part 3	46
Lesson with <code>swirl()</code> : Hierarchical Clustering	46
Lesson 7: K-Means Clustering & Dimension Reduction	47
K-Means Clustering (Part 1)	47
K-Means Clustering (Part 2)	47
Lesson with <code>swirl()</code> : K Means Clustering	47
Dimension Reduction (Part 1)	47
Dimension Reduction (Part 2)	47
Dimension Reduction (Part 3)	47
Lesson with <code>swirl()</code> : Dimension Reduction	47
Lesson with <code>swirl()</code> : Clustering Example	47
Lesson 8: Working with Color in R Plots	47
Part 1	47

Part 2	47
Part 3	47
Part 4	47
Quiz 3 Scribbles	47
Case Studies	48
Clustering Case Study	48
Air Pollution Case Study	48
Lesson with <code>swirl()</code> : Case Study	48
Quiz 4 Scribbles	48
Course Project 2	48

Intro

- Slides and data for this course may be found at [github](#)

Instructor's Note

This course covers the essential exploratory techniques for summarizing data. These techniques are typically applied before formal modeling commences and can help inform the development of more complex statistical models. Exploratory techniques are also important for eliminating or sharpening potential hypotheses about the world that can be addressed by the data. We will cover in detail the plotting systems in R as well as some of the basic principles of constructing data graphics. We will also cover some of the common multivariate statistical techniques used to visualize high-dimensional data.

All the best,

Roger Peng

Introduction

- EDA allows you to develop a rough idea of what your data look like and what kinds of questions might be answered by them.
- EDA is often the “fun part” of data analysis, where you get to play around with the data and explore.
- These techniques for summarizing data are typically applied before formal modeling commences and can help inform the development of more complex statistical models.

Exploratory Data Analysis with R Book

- **Exploratory Data Analysis with R**

The Art of Data Science

- **The Art of Data Science eBook**
- **The Art of Data Science printed version**

Installing R on...

Windows

- Just go to **the cran site** and install the Windows version.
+ For an optimal experience, back up all of data onto a usb, then install your preferred version of Linux (I use Fedora) and install the Linux version instead.

Mac

- Just go to **the cran site** and install the Mac version.
+ If you don't have enough money to buy a Mac install Linux instead, it's open-source, meaning it's free!

Installing R Studio (Mac)

- Install from **the RStudio website** *after* you have R installed.

Setting Your Working Directory on...

Windows, Mac & Linux

- Your working directory is where R will look for all the files it reads and where all the files it writes will go
- `getwd()` will display your current working directory
- `dir()` will display all files in your wd
- `setwd(param)` will set your working directory to the character string that is represented by `param`
- `source("myFunction.R")` will load in `myFunction` script from wd and any functions that are within it.


Lesson 1: Graphs

Principles of Analytic Graphics

- Some general rules to follow when building analytic graphics from data to tell the story the data hold.
- Principles:

1) Show Comparisons


- Evidence for a hypothesis is always *relative* to another competing hypothesis
- Always ask “Compared to What?”
- For example a box plot of **Change in symptom-free days** in children with asthma when an **Air Cleaner** is installed in their quarters should be shown in *comparison* to a control group



Reference: Butz AM, *et al.*, *JAMA Pediatrics*, 2011.

2) Show Causality, Mechanism, or Explanation

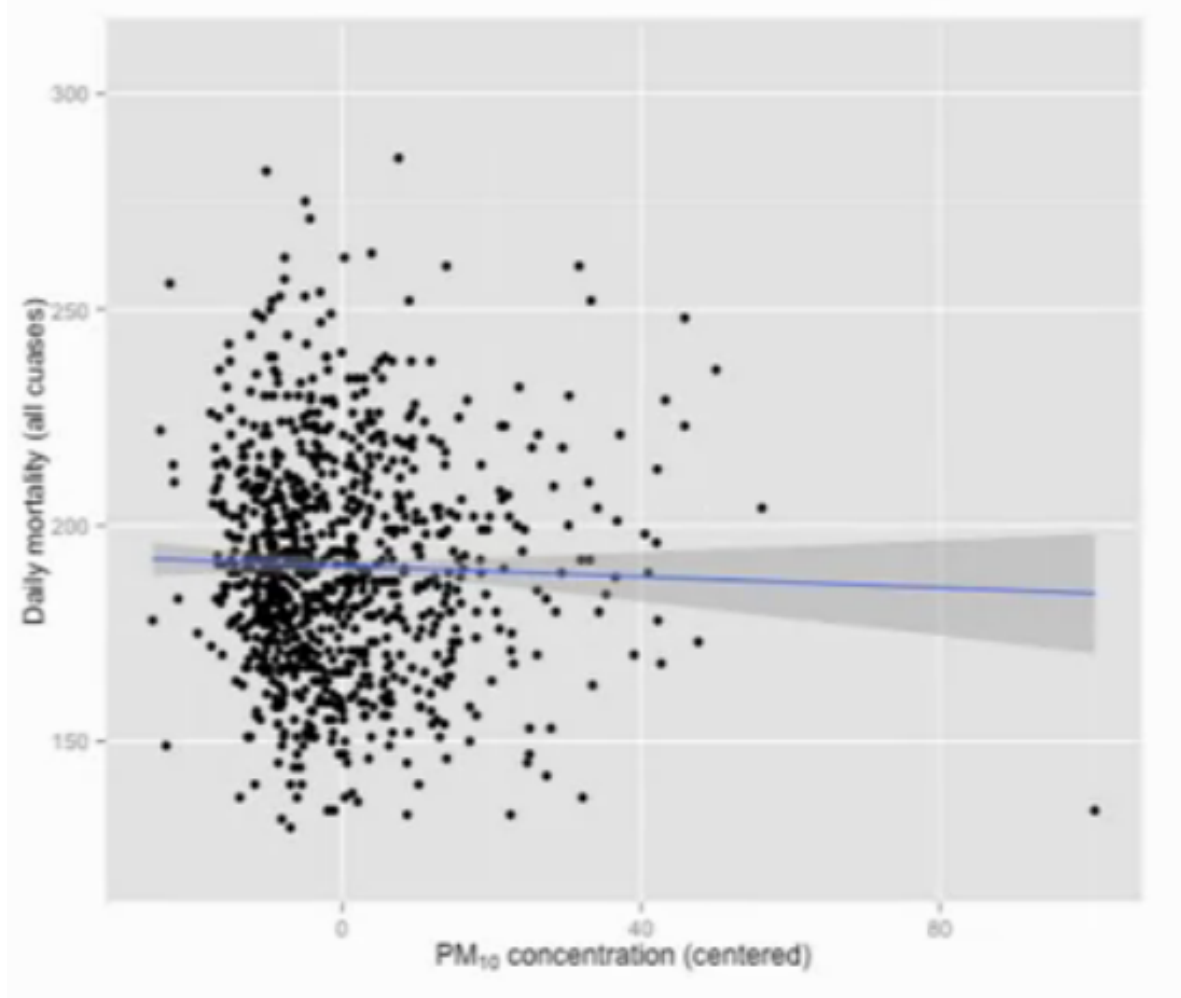
- To show what is going on/how you believe the system is operating and what is the cause for the result you are showing.
- What is your causal framework for thinking about a question
- This only shows a suggestion and indicates where further investigation could go
- In the asthma example you would also want to show the **Change in PM** (Particulate Matter) in the child's home between the **Control** and **Air Cleaner**




3) Show Multivariate Data

- Show as much data on a plot as you reasonably can, it tells a richer story
- The real world is multivariate, so your plots should reflect that

- Need to “escape flatland”
- For example, below is a 2-D plot of **Daily mortality** versus **PM concentration** in NYC, and it shows a slight decrease in mortality as PM increases.




- However, if we plot this across four plots for each season we can see an increase in mortality within each season



4) Integration of Evidence

- Don't let the tool drive the analysis
- Completely integrate words, numbers, images, and diagrams
- Data graphics should make use of many modes of data presentation
- Put lot of information on the plots rather than different places where it may be hard to track down
- The following example shows a plot that has a column for probability that the hospitalizations are different than 0, (the left side is labeling the rows), and the bottom is describing how the experiment was performed

Figure 2. Percentage Change in Emergency Hospital Admissions Rate for Cardiovascular Diseases per a $10\text{-}\mu\text{g}/\text{m}^3$ Increase in Particulate Matter



5) Describe and Document the Evidence

- Use appropriate labels, scales, sources, etc.
- A data graphic should tell a complete story that is also credible

6) Content is King

- If there isn't an interesting story to tell no amount of presentation will make it interesting
- Analytic presentations ultimately stand or fall depending on the quality, relevance, and integrity of their content
- Further Reading - **Edward Tufte's Beautiful Evidence (\$32)**

Lesson with `swirl()`: Principles of Analytic Graphs

- This lesson runs through the 5 principles that were discussed in the above lecture.
- The multivariate plot was an example of **Simpson's paradox, or the Yule-Simpson effect**
- With R, you want to preserve any code you use to generate your data and graphics so that the research can be replicated if necessary.
 - This allows for easy verification or finding bugs in your analysis

Exploratory Graphs

- These are graphs that are made for yourself to look at and explore the data sets you're looking at
- Why do we use graphs in data analysis?
 - To understand data properties
 - To find patterns in data
 - To suggest modeling strategies
 - To “debug” analyses
 - To communicate results
 - Exploratory graphs are for the first four of these reasons
- Characteristics of exploratory graphs
 - They are made quickly (“on the fly”)
 - * Looking through a lot of the variables and different aspects of the data
 - A large number are made
 - * What are the properties, problems, and issues that need followed up
 - Axes/legends are generally cleaned up later
 - Color/size are primarily used for information, rather than presentation

- The following examples will be using data about ambient air pollution in the United States for particle pollution (PM2.5), the “annual mean, averaged over 3 years” cannot exceed 12 micro-grams/cubic meter

```
pollution <- read.csv("./data/avgpm25.csv",
                      colClasses = c("numeric", "character", "factor",
                                     "numeric", "numeric"))
head(pollution)
```

```
##      pm25  fips region longitude latitude
## 1  9.771185 01003   east  -87.74826  30.59278
## 2  9.993817 01027   east  -85.84286  33.26581
## 3 10.688618 01033   east  -87.72596  34.73148
## 4 11.337424 01049   east  -85.79892  34.45913
## 5 12.119764 01055   east  -86.03212  34.01860
## 6 10.827805 01069   east  -85.35039  31.18973
```

- The question we are looking into is: *Do any counties exceed the standard of $12\mu\text{g}/\text{m}^3$?*
 - You always want to have an underlying question in mind, even if it’s kind of vague
- Simple summaries of Data
 - Five-number summary
 - Box plots
 - Histograms
 - Density plot
 - Bar plot

Five-number Summary


- The `summary` function in R gives the 5-number summary as well as the mean

```
summary(pollution$pm25)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      3.383   8.549  10.047   9.836  11.356  18.441
```


Box plots

```
#Lecture used blue but I dislike that display
boxplot(pollution$pm25, col = "hotpink4")
```



- Overlaying a horizontal line to help investigate our question


```
boxplot(pollution$pm25, col = "hotpink4")  
abline(h = 12)
```



Histograms

```
hist(pollution$pm25, col = "green")
```


Histogram of pollution\$pm25



- Including a rug will show detail of the points that causing the plot

```
hist(pollution$pm25, col = "green")  
rug(pollution$pm25)
```


Histogram of pollution\$pm25



- One can also state the number of breaks that are to be in the histogram
 - too big of a number will make too much noise within the histogram
 - too small of a number won't show the shape of the distribution

```
hist(pollution$pm25, col = "green", breaks = 100)
rug(pollution$pm25)
```


Histogram of pollution\$pm25



- Adding a vertical line and the median to the histogram


```
hist(pollution$pm25, col = "green")  
abline(v = 12, lwd = 2) #lwd sets the width of the line  
abline(v = median(pollution$pm25), col = "magenta", lwd = 4)
```

Histogram of pollution\$pm25



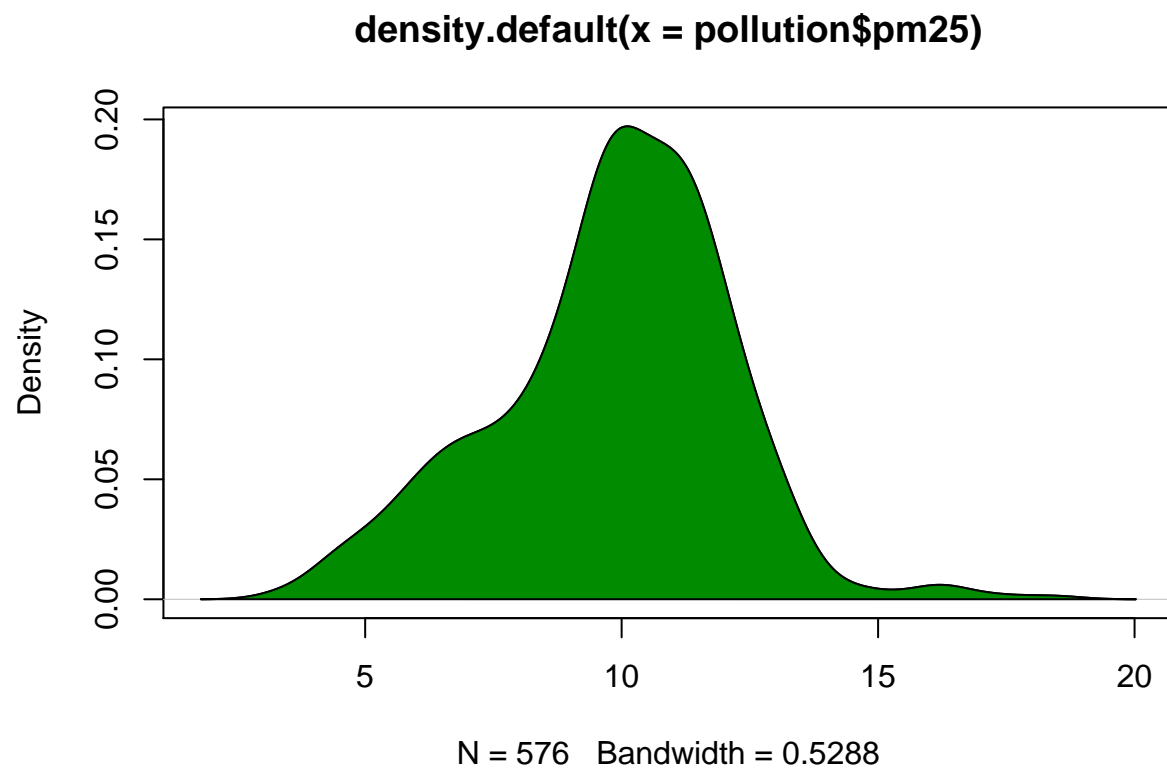
Density plot

```
plot(density(pollution$pm25))
```

- Adding a polygon to fill the area

```
plot(density(pollution$pm25))  
polygon(density(pollution$pm25), col = "green4")
```




Bar plot

- used for comparing categorical variables

```
barplot(table(pollution$region), col = "wheat",  
        main = "Number of Counties in Each Region")
```


Number of Counties in Each Region



- Simple Summaries of Data
 - Two dimensions
 - * Multiple/overlayed 1-D plots (Lattice/ggplot2)
 - * Scatter plots
 - * Smooth scatter plots
 - Greater than 2 dimensions
 - * Overlayed/multiple 2-D plots; coplots
 - * Use color, size, shape to add dimensions
 - * Spinning plots
 - * Actual 3-D plots (not that useful)

Multiple Box plots


```
#Look at pm25 ~(separated by) region  
boxplot(pm25 ~ region, data = pollution, col = "red")
```



Multiple Histograms


```
#mfrow determines the number of: c(row, col)  
#mar is the size of the margins on the c(bottom, left, top, right)  
par(mfrow = c(2, 1), mar = c(4, 4, 2, 1))  
hist(subset(pollution, region == "east")$pm25, col = "green")  
hist(subset(pollution, region == "west")$pm25, col = "green")
```

Histogram of subset(pollution, region == "east")\$pm25



subset(pollution, region == "east")\$pm25


Histogram of subset(pollution, region == "west")\$pm25



subset(pollution, region == "west")\$pm25


Scatter plot

```
with(pollution, plot(latitude, pm25))  
  
#lty = line type  
abline(h = 12, lwd = 2, lty = 2)
```




- Using Color

```
with(pollution, plot(latitude, pm25, col = region))  
abline(h = 12, lwd = 2, lty = 2)
```



- Multiple Scatter plots

```
par(mfrow = c(1,2), mar = c(5, 4, 2, 1))  
with(subset(pollution, region == "west"), plot(latitude, pm25, main = "West"))  
with(subset(pollution, region == "east"), plot(latitude, pm25, main = "East"))
```



- Further Reading
 - **R Graph Gallery**
 - **R Bloggers**

Lesson with `swirl()`: Exploratory Graphs

- Since our brains are very good at seeing patterns, graphs give us a compact way to present data and find or display any pattern that may be present
- We *don't* use exploratory graphs to communicate results
- Exploratory graphs are the “quick and dirty” tool used to point the data scientist in a fruitful direction
- Plot details such as axes, legends, color, and size are cleaned up later to convey more information in an aesthetically pleasing way.

Lesson 2: Plotting


Plotting Systems in R

- Three core plotting systems in R that are useful for achieving various goals

1) Base Plotting System

- Came with original version of R
- Use's the “artist's palette” model of creating graphs, that is it's blank and you pull pieces together
- Start with plot function (or similar)
- Use annotation functions to add/modify (`text`, `lines`, `points`, `axis`)
- Convenient, mirrors how we think of building plots and analyzing data
- Can't go back once the plot is started; can't delete elements
- Difficult to “translate” to others once a new plot has been created (no graphical “language”)
- Plot is just a series of R commands

```
library(datasets)
data(cars)
# Plot speed vs stopping distance
with(cars, plot(speed, dist))
```




2) The Lattice System

- Plots are created with a single function call (`xyplot`, `bwplot`, etc.)
 - Therefore you have to specify a lot of information within the one line
- Most useful for conditioning types of plots: Looking at how y changes with x across levels of z
- Things like margins/spacing set automatically because entire plot is specified at once
- Good for putting many many plots on a screen
- Sometimes awkward to specify an entire plot in a single function call
- Annotation in plot is not especially intuitive
- Use of panel functions and subscripts is difficult to use and requires intense preparation
- Can't "add" anything to a plot after it's created

- The following example shows avg life expectancy in a state versus it's avg income and is divided into four regions


```
library(lattice)
state <- data.frame(state.x77, region = state.region)
xyplot(Life.Exp ~ Income | region, data = state, layout = c(4,1))
```



3) The ggplot2 System

- From “The Grammar of Graphics” by Hadley Wickham (You read this and have it printed out)
- Splits the difference between base and lattice in a number of ways
- Automatically deals with spacings, text, titles but also allows you to annotate by “adding” to a plot
- Superficial similarity to lattice but generally easier/more intuitive to use
- Default mode makes many choices for you (but you can still customize to your heart’s content)
- The following example plots the size of an engine of a car vs the highway mileage of the car

```
library(ggplot2)
data(mpg)
qplot(displ, hwy, data = mpg)
```



- Can't mix and match the different systems because it'll "confuse" the plotting system

Lesson with `swirl()`: Plotting Systems

- `xyplot(Life.Exp ~ Income | region, data = state, layout = c(4,1))`
 - Plots `Life.Exp` (against)~ `Income` (for each)| `region`
 - Data allows us to not have to use `$` in the formula param
 - layout determines the orientation of the graphs

Base Plotting System

Part 1: * The core plotting and graphics engine in R is encapsulated in the following packages:
+ *graphics*: contains plotting functions for the “base” graphing systems, including `plot`, `hist`, `boxplot` and many others. + *grDevices*: contains all the code implementing the various graphics devices, including X11, PDF, PostScript, PNG, etc.

- The lattice plotting system is implemented using the following packages:
 - *lattice*: contains code for producing Trellis graphics, which are independent of the “base” graphics system; includes functions like `xyplot`, `bwplot`, `levelplot`

- *grid*: implements a different graphing system independent of the “base” system; the *lattice* package builds on top of *grid*; we seldom call functions from the *grid* package directly
- When making a plot one must first make a few considerations (not necessarily in this order):
 - Where will the plot be made? On the screen? In a file?
 - How will the plot be used?
 - * Is the plot for viewing temporarily on the screen?
 - * Will it be presented in a web browser?
 - * Will it eventually end up in a paper that might be printed?
 - * Are you using it in a presentation?
 - Is there a large amount of data going into the plot? Or is it just a few points?
 - Do you need to be able to dynamically re-size the graphic?
 - What graphics system will you use: base, lattice, or **ggplot2**? These systems generally cannot be mixed.
 - * Base graphics are usually constructed piecemeal, with each aspect of the plot handled separately through a series of function calls; this is often conceptually simpler and allows plotting to mirror the thought process
 - * Lattice graphics are usually created in a single function call, so all of the graphic’s parameters have to be specified at once; specifying everything at once allows R to automatically calculate the necessary spacings and font sizes
 - * ggplot2 combines concepts from both base and lattice graphics but uses an independent implementation


This lecture focuses on the base plotting system and creating graphics on the screen device only

- * Base graphics are used most commonly and are a very powerful system for creating 2-D graphics
- * There are two phases to creating a base plot + Initializing a new plot + Annotating (adding to) an existing plot
- * Calling `plot(x, y)` or `hist(x)` will launch a graphics device (if one is not already open) and draw a new plot on the device
- * If the arguments to `plot` are not of some special class, then the *default* method for `plot` is called; this function has *many* arguments, letting you set the title, x axis label, y axis label, etc.
- * The base graphics system has *many* parameters that can be set and tweaked; these parameters are documented in `?par`; it wouldn’t hurt to try to memorize this help page

Base Histogram


```
library(datasets)
hist(airquality$Ozone) ## Draw a new plot
```

Histogram of airquality\$Ozone




Base Scatter plot

```
library(datasets)
with(airquality, plot(Wind, Ozone))
```



Base Box plot

```
library(datasets)
airquality <- transform(airquality, Month = factor(Month))
boxplot(Ozone ~ Month, airquality, xlab = "Month", ylab = "Ozone (ppb)")
```



Some Important Base Graphics Parameters

- `pch`: the plotting symbol (default is an open circle)
- `lty`: the line type (default is a solid line), can be dashed, dotted, etc.
- `lwd`: the line width, specified as an integer multiple
 - Thicker lines for presentations as people further back may have trouble seeing it
 - Thinner lines are ok for reports as people can view the plots closer
- `col`: the plotting color, specified as a number, string, or hex code; the `colors()` function gives you a vector of colors by name. I also installed a helpful pdf in the main R folder (`home/phiprime/Documents/Education/R`)
- `xlab`: character string for the x-axis label
- `ylab`: character string for the y-axis label

The `par()` function is used to specify *global* graphics parameters that affect all plots in an R session. These parameters can be overridden when specified as arguments to specific plotting functions. *

`las`: the orientation of the axis labels on the plot

* `bg`: the background color

- * **mar**: the margin size
- * **oma**: the outer margin size (default is 0 for all sides)
- * **mfrow**: number of plots per row, column (plots are filled row-wise)
- * **mfc**: number of plots per row, column (plots are filled column-wise)

- Let's look at some of these defaults:

```
par("lty")

## [1] "solid"

par("col")

## [1] "black"

par("pch")

## [1] 1

par("bg")

## [1] "transparent"

par("mar") #Unit is "lines of text"

## [1] 5.1 4.1 4.1 2.1

par("mfrow")

## [1] 1 1
```


Base Plotting Functions

- **plot**: make a scatter plot, or other type of plot depending on the class of the object being plotted
- **lines**: add lines to a plot, given a vector x values and a corresponding vector of y values (or a 2-column matrix); this function just connects the dots
- **points**: add points to a plot
- **text**: add text labels to a plot using specified x, y coordinates (Inside plot)
- **title**: add annotations to x, y axis labels, title, subtitle, outer margin (Outside plot)
- **mtext**: add arbitrary text to the margins (inner or outer) of the plot
- **axis**: adding axis ticks/labels

Some examples:

Adding title


```
library(datasets)
with(airquality, plot(Wind, Ozone))
title(main = "Ozone and Wind in New York City") ## Add a title
```



Sub-setting some points

```
with(airquality, plot(Wind, Ozone, main = "Ozone and Wind in New York City"))
with(subset(airquality, Month = 5), points(Wind, Ozone, col = "green"))
```


Ozone and Wind in New York City



`type = "n"` will set up the plot but not actually plot any points

```
with(airquality, plot(Wind, Ozone, main = "Ozone and Wind in New York City",
                      type = "n"))
with(subset(airquality, Month == 5), points(Wind, Ozone, col = "green"))
with(subset(airquality, Month != 5), points(Wind, Ozone, col = "red"))
legend("topright", pch = 1, col = c("blue", "red"),
      legend = c("May", "Other Months"))
```


Ozone and Wind in New York City



regression line:



```
with(airquality, plot(Wind, Ozone, main = "Ozone and Wind in New York City",  
                      pch = 20))  
model <- lm(Ozone ~ Wind, airquality)  
abline(model, lwd = 2)
```

Ozone and Wind in New York City



Multiple Base Plots


```
par(mfrow = c(1,2))
with(airquality, {
  plot(Wind, Ozone, main = "Ozone and Wind")
  plot(Solar.R, Ozone, main = "Ozone and Solar Radiation")
})
```



Adding overall plot title:

```
par(mfrow = c(1,3), mar = c(4, 4, 2, 1), oma = c(0, 0, 2, 0))
with(airquality, {
  plot(Wind, Ozone, main = "Ozone and Wind")
  plot(Solar.R, Ozone, main = "Ozone and Solar Radiation")
  plot(Temp, Ozone, main = "Ozone and Temperature")
  mtext("Ozone and Weather in New York City", outer = TRUE)
})
```

Ozone and Weather in New York City




Summary

- Plots in the base plotting system are created by calling successive R functions to “build up” a plot
 - This often results in many lines of code for a plot
- Plotting occurs in two stages:
 - Creation of a plot
 - Annotation of a plot (adding lines, points, text, legends)
- The base plotting system is very flexible and offers a high degree of control over plotting

Base Plotting Demonstration


```
phi <- (1+sqrt(5))/2
x <- rnorm(100)
y <- rnorm(100)
plot(x,y, pch = 4)
text("j", font = 5, x = -phi, y = phi, offset = 0)
```



* Executing `example(Points)` will give a number of demos


```
x <- rnorm(1000)
y <- rnorm(1000)
plot(x, y, pch = ".", xlab = "Weight", ylab = "Height")
title("Scatter plot")
legend("topleft", legend = "Data", pch = ".")
fit <- lm(y ~ x)
abline(fit, lwd = 2, col = "green")
```


Scatter plot



Plotting multiple plots together


```
x <- rnorm(100)
y <- rnorm(100)
z <- rpois(100, 2)
par(mfrow = c(1, 2), mar = c(4, 4, 2, 2))
plot(x,y)
plot(x,z)
```



plotting categorical data

```
x <- rnorm(100)
y <- x + rnorm(100)
g <- gl(2, 50, labels = c("Male", "Female"))
plot(x, y, type = "n")
points(x[g == "Male"], y[g == "Male"], col = "green", pch = 17)
points(x[g == "Female"], y[g == "Female"], col = "red", pch = 15)
legend("topleft", pch = c(17, 15), col = c("green", "red"),
      legend = c("Male", "Female"))


#Adding subsetted regression lines
mfit <- lm(x[g == "Male"] ~ y[g == "Male"])
ffit <- lm(x[g == "Female"] ~ y[g == "Female"])
abline(mfit, col = "green", lwd = 3)
abline(ffit, col = "red", lwd = 3)
```



Lesson with `swirl()`: Base Plotting System

```
boxplot(Ozone~Month, airquality, xlab = "Month", ylab = "Ozone (ppb)",  
        col.axis = "blue", col.lab = "red")  
title("Ozone and Wind in New York City")
```

Ozone and Wind in New York City



```
names(par()) ## Lists parameters
```

```
## [1] "xlog"      "ylog"      "adj"       "ann"       "ask"       "bg"
## [7] "bty"      "cex"       "cex.axis"  "cex.lab"   "cex.main"  "cex.sub"
## [13] "cin"      "col"       "col.axis"  "col.lab"   "col.main"  "col.sub"
## [19] "cra"      "crt"       "csi"       "cxy"       "din"       "err"
## [25] "family"   "fg"        "fig"       "fin"       "font"      "font.axis"
## [31] "font.lab" "font.main" "font.sub"  "lab"       "las"       "lend"
## [37] "lheight"  "ljoin"     "lmitre"    "lty"       "lwd"       "mai"
## [43] "mar"      "mex"       "mfcoll"    "mfg"       "mfrow"     "mgp"
## [49] "mkh"      "new"       "oma"       "omd"       "omi"       "page"
## [55] "pch"      "pin"       "plt"       "ps"        "pty"       "smo"
## [61] "srt"      "tck"       "tcl"       "usr"       "xaxp"      "xaxs"
## [67] "xaxt"     "xpd"       "yaxp"      "yaxs"     "yaxt"     "ylbias"
```

Reminder to commit to GitHub (Delete this line AFTER the commit)

Lesson 3: Graphics Devices

Graphics Devices in R

Lesson with `swirl()`: Graphics Devices in R

Quiz 1 Scribbles

Reminder to commit to GitHub (Delete this line AFTER the commit)

Course Project 1

Reminder to commit to GitHub (Delete this line AFTER the commit)

Lesson 4: Lattice Plotting

Lattice Plotting System (part 1)

Lattice Plotting System (part 2)

Lesson with `swirl()`: Lattice Plotting System

Lesson with `swirl()`: Working with Colors

Reminder to commit to GitHub (Delete this line AFTER the commit)

Lesson 5: ggplot2 <3

Part 1

Part 2

Lesson with `swirl()`: GGPlot2 Part 1

Part 3

Part 4

Lesson with `swirl()`: GGPlot2 Part 2

Part 5

Lesson with `swirl()`: GGPlot2 Extras

Quiz 2 Scribbles

Reminder to commit to GitHub (Delete this line AFTER the commit)

Lesson 6: Hierarchical Clustering

Part 1

Part 2

Part 3

Lesson with `swirl()`: Hierarchical Clustering

Reminder to commit to GitHub (Delete this line AFTER the commit)

Lesson 7: K-Means Clustering & Dimension Reduction

K-Means Clustering (Part 1)

K-Means Clustering (Part 2)

Lesson with `swirl()`: K Means Clustering

Dimension Reduction (Part 1)

Dimension Reduction (Part 2)

Dimension Reduction (Part 3)

Lesson with `swirl()`: Dimension Reduction

Lesson with `swirl()`: Clustering Example

Reminder to commit to GitHub (Delete this line AFTER the commit)

Lesson 8: Working with Color in R Plots

Part 1

Part 2

Part 3

Part 4

Quiz 3 Scribbles

Reminder to commit to GitHub (Delete this line AFTER the commit)

Case Studies

Clustering Case Study

Air Pollution Case Study

Lesson with `swirl()`: Case Study

Quiz 4 Scribbles

Reminder to commit to GitHub (Delete this line AFTER the commit)

Course Project 2

Reminder to commit to GitHub (Delete this line BEFORE the commit)